

Minseok Choi

POSTDOCTORAL FELLOW, KOREA ADVANCED INSTITUTE OF SCIENCE AND TECHNOLOGY (KAIST)

Seoul, South Korea

☎ (+82) 10-7942-9291 | ✉ minseok.choi@kaist.ac.kr | 🏠 brightjade.github.io | 🌐 brightjade | 💼 brightjade | 🐦 @minseok_choi | 🎓 Minseok Choi

Researcher working on trustworthy AI and self-improving agents. My work spans AI safety, model unlearning, red-teaming, and robust agentic systems for real-world applications, with publications at ICLR, ACL, EMNLP, and COLM.

Education

Korea Advanced Institute of Science and Technology (KAIST)

Daejeon, South Korea

INTEGRATED M.S. AND PH.D. IN ARTIFICIAL INTELLIGENCE

Mar. 2020 - Feb. 2026

- GPA: 4.10/4.3, Supervisor: Jaegul Choo
- Thesis: Internal and External Safety Strategies for Robust Large Language Models

University of Illinois at Urbana-Champaign

Urbana, IL

B.S. IN COMPUTER SCIENCE

Aug. 2017 - May. 2019

- GPA: 3.78/4.0, Dean's List: Fall 2017 - Spring 2019

Illinois Institute of Technology

Chicago, IL

B.S. IN COMPUTER SCIENCE

Aug. 2015 - May. 2017

- GPA: 3.89/4.0, Dean's List: Fall 2015 - Spring 2017, International Undergraduate Student Merit Scholarship

Experience

KAIST AI, Data and Visual Analytics Lab (DAVIAN)

Seongnam, South Korea

POSTDOCTORAL FELLOW (SUPERVISOR: JAEGUL CHOO)

Mar. 2026 - Present

- Developing safety methods for self-improving LLM agents.
- InnoCORE Postdoctoral Fellow; concurrently fulfilling national service as Technical Research Personnel through Feb. 2027.

Krafton Inc.

Gangnam, South Korea

APPLIED RESEARCH INTERN (SUPERVISOR: KANGWOOK LEE)

May. 2025 - Feb. 2026

- Developed AI co-playable character (CPC) for PUBG Ally using LLM-based agent systems.
- Contributed to Terminus-KIRA, a coding agent that achieved 74.8% on Terminal-Bench and ranked #2 on the leaderboard (#1 with Opus 4.6).

Naver Webtoon Ltd.

Seongnam, South Korea

RESEARCH INTERN (SUPERVISOR: DAESIK KIM)

Jan. 2022 - Jun. 2022

- Researched keyphrase extraction and generation methods for webtoon content, contributing to SimCKP (EMNLP'23).

Publications

* indicates equal contribution.

PREPRINTS

- P2** **Minseok Choi**,* Seungbin Yang,* Dongjin Kim,* Subin Kim, Jungmin Son, Yunseung Lee, Jaegul Choo, and Youngjun Kwak. "Membrane: A Self-Evolving Contrastive Safety Memory for LLM Agent Defense," Preprint, 2026.
- P1** **Minseok Choi**, ChaeHun Park, Dohyun Lee, and Jaegul Choo. "Breaking Chains: Unraveling the Links in Multi-Hop Knowledge Unlearning," Preprint, 2025.

PEER-REVIEWED CONFERENCE/JOURNAL PAPERS

- C9** **Minseok Choi**,* Dongjin Kim,* Seungbin Yang,* Subin Kim, Youngjun Kwak, Juyoung Oh, Jaegul Choo, and Jungmin Son. "Exp-Guard: LLM Content Moderation in Specialized Domains," International Conference on Learning Representations (ICLR'26), 2026.
- C8** **Minseok Choi**, Daniel Rim, Dohyun Lee, and Jaegul Choo. "Opt-Out: Investigating Entity-Level Unlearning for Large Language Models via Optimal Transport," Annual Meeting of the Association for Computational Linguistics (ACL'25), 2025.
- C7** **Minseok Choi**, Kyunghyun Min, and Jaegul Choo. "Cross-Lingual Unlearning of Selective Knowledge in Multilingual Language Models," Conference on Empirical Methods in Natural Language Processing (EMNLP'24), Findings of EMNLP, 2024.

C6

ChaeHun Park, **Minseok Choi**, Dohyun Lee, and Jaegul Choo. “PairEval: Open-domain Dialogue Evaluation with Pairwise Comparison,” Conference on Language Modeling (**COLM’24**), 2024.

C5

Dohyun Lee*, Daniel Rim*, **Minseok Choi**, and Jaegul Choo. “Protecting Privacy Through Approximating Optimal Parameters for Sequence Unlearning in Language Models,” Annual Meeting of the Association for Computational Linguistics (**ACL’24**), Findings of ACL, 2024.

C4

Minseok Choi, Chaeheon Gwak, Seho Kim, Si Hyeong Kim, and Jaegul Choo. “SimCKP: Contrastive Learning of Keyphrase Representations,” Conference on Empirical Methods in Natural Language Processing (**EMNLP’23**), Findings of EMNLP, 2023.

C3

Minseok Choi, Hyesu Lim, and Jaegul Choo. “PRISM: Enhancing Low-Resource Document-Level Relation Extraction with Relation-Aware Score Calibration,” International Joint Conference on Natural Language Processing and the Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics (**IJCNLP-AAACL’23**), Findings of IJCNLP-AAACL, 2023.

C2

Soyoung Yang, **Minseok Choi**, Youngwoo Cho, and Jaegul Choo. “HistRED: A Historical Document-level Relation Extraction Dataset,” Annual Meeting of the Association for Computational Linguistics (**ACL’23**), 2023.

C1

Hojun Cho, Dohee Kim, Seungwoo Ryu, ChaeHun Park, Hyungjong Noh, Jeong-in Hwang, **Minseok Choi**, Edward Choi, and Jaegul Choo. “Rethinking Style Transformer by Energy-based Interpretation: Adversarial Unsupervised Style Transfer using Pretrained Model,” Conference on Empirical Methods in Natural Language Processing (**EMNLP’22**), 2022.

J1

Minseok Choi, Jaegul Choo, and Minsu Ha. “Development and Application of Web-based Machine Learning Program for Automated Assessment Model Generation,” Brain, Digital, & Learning, Vol. 12, No. 4, pp. 567-578, 2022.

Core Expertise

Research Areas	Trustworthy AI, Model Unlearning, Self-Improving Agents
Applied Skills	Agent Design & Orchestration, LLM Fine-Tuning, Red-Teaming, Safety Evaluation
Infrastructure	Python, PyTorch, HF Transformers/PEFT, vLLM, DeepSpeed, Weights & Biases
Languages	Bilingual: English (professional proficiency; lived in US 2008-2019) and Korean (native)

Teaching & Outreach

SK Advanced Practitioner Exam for NLP	Seoul, South Korea
EXAMINER	Mar.-Sep. 2024 & Feb.-May. 2025
SK AI Data Academy for NLP	Seoul, South Korea
TEACHING ASSISTANT	Jan. 2023 - Mar. 2023
Seoul Machine Learning Boot Camp	Seoul, South Korea
TEACHING ASSISTANT	Jun. 2022 - Jul. 2022
Startup AI Boot Camp (YearDream)	Seoul, South Korea
TEACHING ASSISTANT	Sep.-Dec. 2021 & Apr.-May. 2024

Professional Service

2023-2026	Reviewer , Annual Meeting of the Association for Computational Linguistics (ACL)
2023-2026	Reviewer , Conference on Empirical Methods in Natural Language Processing (EMNLP)
2023-2026	Reviewer , Annual Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)
2026	Reviewer , Conference on Neural Information Processing Systems (NeurIPS)
2026	Reviewer , AAAI Conference on Artificial Intelligence (AAAI)
2025	Reviewer , International Conference on Computational Linguistics (COLING)